

- 
- Review
  - Dynamic Memory Allocation
    - Designing the `malloc()` Function
    - A Simple Implementation of `malloc()`
    - A Real-World Implementation of `malloc()`
  - Using `malloc()`
  - Garbage Collection

# Review: C standard library

---

- **I/O functions:** `fopen()`, `freopen()`, `fflush()`, `remove()`, `rename()`, `tmpfile()`, `tmpnam()`, `fread()`, `fwrite()`, `fseek()`, `ftell()`, `rewind()`, `clearerr()`, `feof()`, `ferror()`
- **Character testing functions:** `isalpha()`, `isdigit()`, `isalnum()`, `isctrl()`, `islower()`, `isprint()`, `ispunct()`, `isspace()`, `isupper()`
- **Memory functions:** `memcpy()`, `memmove()`, `memcmp()`, `memset()`

# Review: C standard library

---

- **Conversion functions:** `atoi()`, `atol()`, `atof()`, `strtol()`, `strtoul()`, `strtod()`
- **Utility functions:** `rand()`, `srand()`, `abort()`, `exit()`, `atexit()`, `system()`, `bsearch()`, `qsort()`
- **Diagnostics:** `assert()` function, `__FILE__`, `__LINE__` macros

# Review: C standard library

---

- Variable argument lists:
  - Declaration with ... for variable argument list (may be of any type):  
`int printf(const char * fmt, ...);`
  - Access using data structure `va_list ap`, initialized using `va_start()`, accessed using `va_arg()`, destroyed at end using `va_end()`
- Time functions: `clock()`, `time()`, `difftime()`, `mktime()`, `asctime()`, `localtime()`, `ctime()`, `strftime()`

- 
- Review
  - **Dynamic Memory Allocation**
    - Designing the `malloc()` Function
    - A Simple Implementation of `malloc()`
    - A Real-World Implementation of `malloc()`
  - Using `malloc()`
  - Garbage Collection

# Dynamic memory allocation

---

- Memory allocated during runtime
- Request to map memory using `mmap ( )` function (in `<sys/mman.h>`)
- Virtual memory can be returned to OS using `munmap ( )`
- Virtual memory either backed by a file/device or by *demand-zero* memory:
  - all bits initialized to zero
  - not stored on disk
  - used for stack, heap, uninitialized (at compile time) globals

# Mapping memory

---

- Mapping memory:

```
void *mmap(void *start, size_t length, int prot,  
           int flags, int fd, off_t offset);
```

- asks OS to map virtual memory of specified length, using specified physical memory (file or demand-zero)
  - `fd` is file descriptor (integer referring to a file, not a file stream) for physical memory (i.e. file) to load into memory
  - for demand-zero, including the heap, use `MMAP_ANON` flag
  - `start` – suggested starting address of mapped memory, usually `NULL`
- Unmap memory:  

```
int munmap(void *start, size_t length);
```

# The heap

---

- Heap – private section of virtual memory (demand-zero) used for dynamic allocation
- Starts empty, zero-sized
- `brk` – OS pointer to top of heap, moves upwards as heap grows
- To resize heap, can use `sbrk()` function:  
`void *sbrk(int inc); /* returns old value of brk_ptr */`
- Functions like `malloc()` and `new` (in C++) manage heap, mapping memory as needed
- Dynamic memory allocators divide heap into blocks



# Requirements

---

- Must be able to allocate, free memory in any order
- Auxiliary data structure must be on heap
- Allocated memory cannot be moved
- Attempt to minimize fragmentation

# Fragmentation

---

- Two types – internal and external
- Internal – block size larger than allocated variable in block
- External – free blocks spread out on heap
- Minimize external fragmentation by preferring fewer larger free blocks

# Design choices

---

- Data structure to track blocks
- Algorithm for positioning a new allocation
- Splitting/joining free blocks

# Tracking blocks

---

- Implicit free list: no data structure required
- Explicit free list: heap divided into fixed-size blocks; maintain a linked list of free blocks
  - allocating memory: remove allocated block from list
  - freeing memory: add block back to free list
- Linked list iteration in linear time
- Segregated free list: multiple linked lists for blocks of different sizes
- Explicit lists stored within blocks (pointers in payload section of free blocks)

## Positioning allocations

---

- Block must be large enough for allocation
- First fit: start at beginning of list, use first block
- Next fit: start at end of last search, use next block
- Best fit: examines entire free list, uses smallest block
- First fit and next fit can fragment beginning of heap, but relatively fast
- Best fit can have best memory utilization, but at cost of examining entire list

# Splitting and joining blocks

---

- At allocation, can use entire free block, or part of it, splitting the block in two
- Splitting reduces internal fragmentation, but more complicated to implement
- Similarly, can join adjacent free blocks during (or after) freeing to reduce external fragmentation
- To join (coalesce) blocks, need to know address of adjacent blocks
- Footer with pointer to head of block – enable successive block to find address of previous block

## Explicit free list

---

- Maintain pointer to head, tail of free list (not in address order)
- When freeing, add free block to end of list; set pointer to next, previous block in free list at beginning of payload section of block
- When allocating, iterate through free list, remove from list when allocating block
- For segregated free lists, allocator maintains array of lists for different sized free blocks

## `malloc()` for the real world

---

- Used in GNU libc version of `malloc()`
- Details have changed, but nice general discussion can be found at  
<https://gee.cs.oswego.edu/dl/html/malloc.html>
- Chunks implemented as in segregated free list, with pointers to previous/next chunks in free list in payload of free blocks
- Lists segregated into bins according to size; bin sizes spaced logarithmically
- Placement done in best-fit order
- Deferred coalescing and splitting performed to minimize overhead



- 
- Review
  - Dynamic Memory Allocation
    - Designing the `malloc()` Function
    - A Simple Implementation of `malloc()`
    - A Real-World Implementation of `malloc()`
  - **Using `malloc()`**
  - Garbage Collection

## Using `malloc()`

---

- Minimize overhead – use fewer, larger allocations
- Minimize fragmentation – reuse memory allocations as much as possible
- Growing memory – using `realloc()` can reduce fragmentation
- Repeated allocation and freeing of variables can lead to poor performance from unnecessary splitting/coalescing (depending on implementation of `malloc()`)

- 
- Review
  - Dynamic Memory Allocation
    - Designing the `malloc()` Function
    - A Simple Implementation of `malloc()`
    - A Real-World Implementation of `malloc()`
  - Using `malloc()`
  - Garbage Collection

# Garbage collection

---

- C implements no garbage collector
- Memory not freed remains in virtual memory until program terminates
- Other languages like Java implement garbage collectors to free unreferenced memory
- When is memory unreferenced?

# Garbage collection

---

- C implements no garbage collector
- Memory not freed remains in virtual memory until program terminates
- Other languages like Java implement garbage collectors to free unreferenced memory
- When is memory unreferenced?
  - Pointer(s) to memory no longer exist
  - Tricky when pointers on heap or references are circular (think of circular linked lists)
  - Pointers can be masked as data in memory; garbage collector may free data that is still referenced (or not free unreferenced data)

# Garbage collection and memory allocation

---

- Program relies on garbage collector to free memory
- Garbage collector calls `free()`
- `malloc()` may call garbage collector if memory allocation above a threshold

# Mark and sweep garbage collector

---

- Simple tracing garbage collector
- Starts with list of known in-use memory (e.g. the stack)
- Mark: trace all pointers, marking data on the heap as it goes
- Sweep: traverse entire heap, freeing unmarked data
- Requires two complete traversals of memory, takes a lot of time

# Copying garbage collector

---

- Uses a duplicate heap; copies live objects during traversal to the duplicate heap (the *to-space*)
- Updates pointers to point to new object locations in duplicate heap
- After copying phase, entire old heap (the *from-space*) is freed
- Code can only use half the heap



# Summary

---

Topics covered:

- Dynamic memory allocation
  - the heap
  - designing a memory allocator
  - a real world allocator
- Using `malloc()`  
↘
- Garbage collection
  - mark-and-sweep collector
  - copying collector